Lateral Tree-of-Thoughts Surpasses ToT by Incorporating Logically-Consistent, Low-Utility Candidates

Abhinav Madahar [1]

Abstract

Modern deployments increasingly budget *large test-time compute*—thousands of tokens or many node expansions—to improve reliability. When *structured search* (e.g., ToT/MCTS-style controllers) is run under such frontier-like conditions, two effects intensify: *breadth saturation*, where additional samples at a node mostly yield near-duplicates so width stops growing; and *myopia*, where early, noisy utility undervalues branches whose payoff appears only after a few more steps, so they are pruned too soon. We introduce **Lateral Tree-of-Thoughts (LToT)**, a search-time controller that explicitly separates the frontier into *mainlines—high-utility* candidates used for exploitation—and *laterals—logically consistent, initially low-utility* candidates that merit short, cheap probes before judgment. LToT explores laterals via *Lateral Racing with Short-Circuit (LR–SC)*, a budgeted race that spreads tiny probes across a very wide lateral set, culls aggressively, and immediately promotes a branch once it demonstrably clears the mainline bar; mainlines are kept intentionally narrow so surplus compute is invested where width is cheap. This turns large budgets into principled diversity while preserving promotion discipline. Let N_0 denote the *initial lateral width* (the number of laterals admitted to the race) and $\eta > 1$ the *culling factor* between rungs. We show a *pseudolinear lateral cost* $\Theta(N_0 \log_\eta N_0)$ with logarithmically many rungs. Pseudolinearity matters because it allows lateral width to scale almost linearly in cost, so increasing budgets buy *coverage* rather than redundant deepening. Under *equal compute*, we evaluate on GSM-Hard/Plus, MATH-500, HumanEval, MBPP-lite, and Game-of-24, reporting Success@1/Pass@1, width scaling, *time-to-first-verified* solution, and *false-promotion* rates. Across math, code, and ToT-style puzzles, LToT improves or matches accuracy while *reducing expansions-to-first-hit*, converting surplus test-time compute into *pro-*

ductive breadth without sacrificing selectivity.

1. Introduction

Modern language models (LMs) are increasingly deployed in *compute–rich* inference regimes: users and systems budget thousands of tokens or many node expansions per query in return for reliability. The dominant recipe for using this budget is structured search, most commonly Tree–of–Thoughts (ToT) over partial solutions (Yao et al., 2023a), layered on top of stepwise prompting (Wei et al., 2022; Wang et al., 2022; Kojima et al., 2022).

We argue that inference–time controllers should treat *logically consistent, low–utility* candidates as assets, not waste. The architectural move is to separate *consistency/continuability* from *utility*, and to invest cheap budget into short *predictive continuations* of many consistent branches until a small number of them demonstrate sustained marginal improvement. This idea echoes "lateral thinking" (de Bono, 1967): rather than deepening a few promising threads, maintain wide, low–commitment exploration that can be *promoted* the instant it proves itself. We introduce **Lateral Tree–of–Thoughts (LToT)**, a drop–in search–time controller that operationalizes the thesis above.

The design converts extra tokens into *principled diversity* where it is cheapest: lateral width. A simple cost law shows that LR–SC's lateral spend is $\Theta(N_0 \log_\eta N_0)$ in initial width N_0 (constant cost per rung and $O(\log_\eta N_0)$ rungs), while mainlines would grow exponentially with depth if left uncapped. By allocating on *marginal improvement* rather than level, LR–SC rescues branches that a single noisy utility reading would discard, while width–aware thresholds prevent "lucky spikes" from polluting mainlines. In aggregate, LToT mitigates breadth saturation and depth myopia without inflating compute. We propose LToT, a controller that keeps mainlines narrow and pushes exploration laterally through LR–SC— a successive–halving race with short–circuit promotion, robust order–aware detection, and freeze–thaw survivors. We bind promotion to verifier–aligned outcomes (exact match/tests for math/code) and introduce a dual gate for

QA (plausibility and path consistency), reducing mainline contamination under noisy evaluators. We derive width–aware bars (sub–Gaussian, sub–Gamma, and sub–Weibull variants) and a repeat–to–confirm rule that keep false promotions bounded as lateral width grows; we also handle correlation via an effective width.

We prove a pseudolinear lateral cost law $\Theta(N_0 \log_\eta N_0)$, logarithmic rung depth, and error bounds under mild tail assumptions, and contrast this with exponential growth in uncapped mainlines. Across math (GSM variants, MATH–500), code (HumanEval/MBPP–lite), and a canonical ToT puzzle (Game of 24), LToT improves Success@1/Pass@1 at matched compute over CoT, vanilla ToT, and MCTS with progressive widening (Xie et al., 2024), while lowering false promotions and time–to–first–hit via short–circuiting.

LToT complements inference–time scaling via best–of–n (Chen et al., 2024; Yang et al., 2024) and revising/self–improvement (Madaan et al., 2023), and sits alongside program/tool–aided reasoning (Gao et al., 2022; Chen et al., 2023). Its novelty is not another search heuristic but a *control principle*: separate consistency from utility; allocate on marginal improvement; and convert surplus compute into lateral breadth with guarantees.

2. Motivation

Frontier language models increasingly run in *compute-rich* inference settings: users and systems are willing to spend thousands of tokens (or many node expansions) per query to improve reliability. Yet the dominant search pattern—vanilla Tree-of-Thoughts (ToT)—*under-utilizes* this budget in two systematic ways already visible today and poised to worsen as budgets grow:

1. **Utility saturation (breadth collapse).** After a handful of genuinely distinct high-utility continuations, additional samples at a node mostly yield near-duplicates whose v scores fall just below the pruning threshold. The frontier then remains narrow even when ample budget is available, leaving compute unused.

2. **Myopic pruning (depth myopia).** Early v estimates are noisy and biased toward near-term payoff; logically consistent branches whose payoff is delayed by several steps are pruned as "low-v" even though they could mature into correct solutions. This creates *myopic false negatives*.

Both effects amplify with larger inference budgets: saturation wastes more samples as k grows, and myopic pruning discards more candidates as depth increases. Let k be the number of children sampled per expanded node and let a be the acceptance fraction into the *mainline*. If one does not cap mainline width, the expected number of mainline nodes at depth d scales like $(ak)^d$, so the cost to depth D is $\Theta((ak)^D)$—*exponential in depth*. By contrast, controlling *lateral* width with successive-halving (LR-SC; Sec. 4.2) yields a total lateral exploration cost of $\Theta(N_0 \log_\eta N_0)$ for initial lateral width N_0 and culling factor $\eta > 1$—*pseudolinear in width*. This asymmetry suggests an architectural principle: *keep mainlines narrow to avoid depth explosion and push width into laterals where it is cheap.* Three trends sharpen the pain points above.

As budgets per query rise, multi-round agents, tool use, and safety/verification passes make higher spend acceptable; yet, without a controller that converts extra tokens into *productive breadth*, vanilla ToT quickly saturates and the marginal return of compute collapses. The problem is amplified as more tasks exhibit long horizons: program synthesis, multi-hop reasoning, and formal verification often require several structured steps before v improves, so naive pruning disproportionately excises precisely those candidates that need brief nurturing. Compounding this, practical evaluators v are noisy and non-stationary across depth and domain; a fixed, level-based gate conflates noise with signal, so allocation must track the *marginal value of compute* rather than a single snapshot.

Let a candidate node x have an (unobserved) eventual value $\mu(x)$ if its branch were fully developed. An early evaluator observes $v(x) = \mu(x) - \lambda \Delta(x) + \varepsilon$, where $\Delta(x)$ is the (unknown) remaining steps to payoff, $\lambda > 0$ captures horizon bias, and ε is evaluator noise. When $\Delta(x)$ is moderate, $v(x)$ may fall below the mainline gate despite large $\mu(x)$. A controller that reasons about *improvement after a small investment*—rather than $v(x)$ in isolation—can *defer judgment*, test whether x starts producing high-v descendants quickly, and only then commit budget. The remedy is to decide continuation on *compute-normalized* improvement—the observed slope (and curvature) of v as budget increases—rather than an absolute v at a preselected level. This steers budget toward branches whose incremental yield remains positive and away from ones whose progress has flattened.

Exploration should be wide but short: spread tiny probes across a very large lateral set, cull aggressively, and snap back to exploitation the moment any lateral reaches the *mainline bar*. Keep mainlines narrow via beam or quota caps to avoid the $(ak)^D$ depth blow-up, and re-open exploration when exploitation *plateaus* in compute-normalized progress. Promote only on outcomes validated by a v that is as outcome-aligned as possible so that speciously plausible branches do not contaminate the mainline. As lateral width grows, control multiplicity by deduplicating near-duplicates, using width-aware thresholds, and adding a cheap repeat-to-confirm step so that evidence is not double-counted.

LToT operationalizes these desiderata with two ingredients (see Sec. 4): (i) a *dual-score frontier* that retains logically consistent, low-v *laterals* alongside high-v *mainlines*, deferring judgment on laterals; and (ii) a budgeted racing procedure, *LR-SC*, that allocates tiny probes across a very wide lateral set, culls aggressively, and *promotes* a lateral to the exploitation set the moment its envelope reaches the mainline bar. Theoretical analyses (Sec. 4.4) show that LR-SC keeps lateral cost *pseudolinear in width* ($\Theta(N_0 \log_\eta N_0)$) while mainlines, if left uncapped, are exponential in depth; hence LToT converts surplus compute into principled diversity exactly where it is cheapest.

3. Related Work

A large body of work elicits multi-step reasoning at inference time by prompting language models to externalize intermediate steps. Chain-of-Thought (CoT) (Wei et al., 2022) and Zero-shot CoT (Kojima et al., 2022) demonstrate that free-form rationales can substantially improve performance on math, symbolic, and commonsense tasks. Several variants structure this process: Self-Consistency aggregates multiple CoT samples via voting to reduce variance (Wang et al., 2022); Least-to-Most decomposes problems into sub-questions solved sequentially (Zhou et al., 2022); Plan-and-Solve asks models to sketch a plan before executing it (Wang et al., 2023); and ReAct interleaves short reasoning traces with tool-use actions (Yao et al., 2023b). These methods focus primarily on generating and aggregating linear traces; in contrast, our Lateral Tree-of-Thoughts (LToT) explicitly organizes alternatives in a *tree* while preserving logically consistent but low-utility branches to improve global search coverage. See also Press et al. (2023) on self-questioning for decomposition.

Tree-of-Thoughts (ToT) casts reasoning as a search over partial thoughts with learned/heuristic evaluators (Yao et al., 2023a). Subsequent work generalizes the structure from trees to graphs (Graph-of-Thoughts) (Besta et al., 2024), ensembles multiple trees (Forest-of-Thought) (Bi et al., 2024), and explores "Everything-of-Thoughts" style meta-frameworks (Ding et al., 2023). Efficiency-oriented advances include Dynamic Parallel Tree Search (DPTS), which parallelizes ToT expansions and focuses compute on promising branches (Ding et al., 2025). Our approach is complementary: rather than accelerating a fixed search policy or collapsing branches early, LToT *retains* laterally related, logically consistent candidates that appear locally low-utility, improving the chance of escaping premature pruning and enabling cross-branch re-use of partial deductions.

A parallel line of work iteratively revises solutions. Self-Refine uses the model's own feedback to edit drafts (Madaan et al., 2023); Reflexion stores episodic "verbal reinforcement" to guide future trials (Shinn et al., 2023). "Boosting/Buffer-of-Thoughts" families build and retrieve reusable thought templates or ensembles to improve robustness and cost (Chen et al., 2024; Yang et al., 2024). LToT differs in objective and mechanism: it organizes contemporaneous alternatives in a search tree and deliberately curates *logically consistent, low-scored* branches to maintain breadth, rather than relying on post-hoc reflection or global templates.

Scaling inference-time compute via repeated sampling ("best-of-N") and diverse rationales improves accuracy when paired with selection mechanisms (Cobbe et al., 2021; Wang et al., 2022). Recent studies formalize test-time scaling and its limits, highlighting that simple majority vote and naïve reward models can plateau, while coverage grows with sample budget (Brown et al., 2024–2025). Verifier training and process supervision further enhance selection quality (Lightman et al., 2023; Zhang & coauthors, 2024). LToT contributes a complementary lever: rather than solely increasing samples or verifier strength, it *rebalances* exploration by preserving branches that are logically sound yet temporarily low-utility, improving coverage of the hypothesis space under fixed compute.

Program-Aided Language Models (PAL) and Program-of-Thoughts (PoT) separate symbolic computation from natural-language reasoning by delegating computation to interpreters (Gao et al., 2022; Chen et al., 2023). Such modularization can be combined with structured search: MCTS-style controllers over chains of thought (Xie et al., 2024), and process-reward models with MCTS (OmegaPRM) (Luo et al., 2024). LToT is orthogonal: it can host code-executed checks inside nodes while preserving lateral candidates that pass logical checks but score poorly under short-horizon utilities.

Although CoT often boosts task accuracy, generated rationales may be unfaithful (Turpin et al., 2023; Lanham et al., 2023). Methods to improve faithfulness include "faithful-by-construction" pipelines that deterministically execute symbolic traces (Lyu et al., 2023) and self-verification prompts or analyses (Weng et al., 2023). Our emphasis on *logical consistency* as a retention criterion naturally interacts with these concerns: LToT filters and preserves candidates whose internal derivations satisfy logical checks, even when immediate utility scores are low, aligning exploration pressure with consistency rather than only with myopic reward.

Finally, techniques such as Skeleton-of-Thought prompt models to outline and then expand subparts in parallel, reducing latency while sometimes improving quality (Ning et al., 2024). Orthogonal to latency, LToT targets *exploration completeness*: by laterally preserving logically consistent branches, it trades small additional compute for a

LToT element	Closest prior	What is different here
Predictive continuation	SH/Hyperband levels	Forecasted marginal gain on a branch *envelope*; order-aware bar with confirmation
Width-aware bar + confirm	Heuristics	Explicit $\log(|S_r||\mathcal{M}_r|)$ control; heavy-tail variants; effective width for correlation
Verifier-bound promotion	Budget milestones	Promotion tied to exact tests / EM; dual gate under plausibility
Short-circuit to exploit	Bracket completion	Immediate return upon meeting mainline bar $B_t + \delta$
Freeze–thaw laterals	Freeze–thaw BO	Applied to reasoning traces with cached rung state
Dual-score frontier	—	Distinguishes high-v mainlines vs. high-c, low-v laterals

Table 1: Cross-walk: LToT control rules vs. racing backbones.

disproportionate increase in the chance of reaching globally consistent solutions under bounded budgets.

We adopt a successive-halving (racing) backbone solely to control lateral cost (pseudolinear $\Theta(N_0 \log_\eta N_0)$, logarithmic rungs). The novelty in LToT lies in reasoning-specific *control rules* layered atop this backbone: (i) compute-normalized predictive continuation on branch envelopes (local polynomial forecast); (ii) width-aware thresholds with confirmation to control max-of-many effects; (iii) verifier-aligned promotion (dual-gated under plausibility); (iv) short-circuit to exploitation on success; (v) freeze–thaw of laterals across phases; and (vi) a dual-score frontier separating high-v mainlines from high-c, low-v laterals. Ablations and SH-only baselines (Sec. 5) show that the backbone alone does not yield our accuracy, false-promotion, or latency characteristics.

4. Architecture Design

LToT is a search-time controller for reasoning with language models (LMs) that (i) keeps *mainlines* narrow to avoid exponential blow-up in depth and (ii) makes *lateral* exploration very wide but cheap via a budgeted racing procedure with short-circuit promotion. The controller decides when to exploit mainlines vs. explore laterals, and—during exploration—how to allocate compute across many lateral branches while maintaining guarantees on cost and false promotions.

We reason over a rooted tree (or DAG) of partial traces. Each node x is a partial solution; its children are produced by prompting the LM with x. Two evaluators score nodes:

$v(x) \in \mathbb{R}$ (utility; e.g., answer- or verifier-aligned), $c(x) \in [0, 1]$ (logical consistency / soundness).

We measure compute in either node expansions or tokens and denote cumulative compute by C. For any node x with parent p, we define a *local consistency* score

$$c_{\text{local}}(x) = \lambda_1\, s_{\text{logic}}(x \mid p) + \lambda_2\, s_{\text{syntax}}(x) + \lambda_3\, s_{\text{constraints}}(x),$$
(1)

where s_{logic} is an LM step-checker that validates whether the new line follows from the previous state, s_{syntax} checks

parsability/format (e.g., code compiles, expression parses), and $s_{\text{constraints}}$ encodes simple domain invariants (e.g., no new free variables, signature preserved). If a component is unavailable we reweight the remaining terms proportionally; when s_{logic} carries $\lambda_1 \geq 0.7$ we tighten the promotion gate in Sec. 4.3 by raising the path-consistency threshold by $+0.1$ and requiring one-step re-derivation. We aggregate consistency along a branch i of length h via a robust *path-consistency* score

$$C_{\text{path}}(i, h) = \min\Big\{ \text{Quantile}_q\left(\{c_{\text{local}}(x_j)\}_{j \leq h}\right),\, \bar{c}_{\text{local}}(i, h) \Big\}, \quad q = 0.25,$$
(2)

which is distribution-free and stable for short paths. (A mean−MAD variant appears in App. A.) Each branch i maintains a tiny *micro-beam* of size m_μ leaves at each horizon. We define the *envelope* at horizon h as a smoothed Top-K mean over those leaves,

$$V_i(h) = \text{TopKMean}\big(\{v(\ell)\}_{\ell \in \mathcal{L}_i(h)}; K\big), \quad K = m_\mu,$$
(3)

with Beta smoothing

$$\tilde{V}_i(h) = \frac{K_* V_i(h) + \alpha}{K_* + 2\alpha}, \quad \alpha = 0.5,$$
(4)

where $K_* = K$ for Top-K. Optionally we use a weighted envelope $V_i(h) = \sum_{j=1}^{m_\mu} \omega_{ij}\, v_{ij}$ with clipped-softmax weights $0 \leq \omega_{ij} \leq \omega_{\max}$, $\sum_j \omega_{ij} = 1$; we then set the effective sample size $K_* = K_{\text{eff}} = 1/\sum_j \omega_{ij}^2$ in the smoothing formula. This adapts the shrinkage to how many leaves effectively contribute and stabilizes the continuation statistic. Unless stated otherwise we set $m_\mu = 3$, $K = m_\mu$. At time t the search maintains a frontier $\mathcal{F}_t$ and an *exploitation set* $M_t \subseteq \mathcal{F}_t$ of nodes eligible for *mainline* exploitation. Nodes carry an immutable origin tag in {MAINLINE_ORIGIN, LATERAL_ORIGIN} indicating how they first entered the frontier. We also maintain a *mainline acceptance bar* B_t (e.g., the best-so-far v or a top-k mean with a small margin $\delta > 0$).

Children with high v are admitted to M_t (mainlines). Children with low v but high c enter the *lateral pool* L_t for potential exploration. Intuitively, laterals represent hypotheses that appear unpromising under a myopic utility but are logically coherent and may become valuable after a short lookahead.

For a lateral branch i (rooted at node x_i), let $V_i(h)$ denote a branch *envelope*—e.g., a Top-k mean of v among leaves within horizon h steps from x_i (or within a per-branch micro-beam). We write $C(h)$ for the compute required to reach horizon h_j and define the compute-normalized improvement between horizons $h' < h$ as

$$g_i(h, h') = \frac{V_i(h) - V_i(h')}{C(h) - C(h')}.$$

4

These quantities let us reason about *marginal value of compute*, not just absolute levels.

4.1. Controller overview

LToT alternates between: Expand nodes from M_t while a compute-normalized progress statistic (e.g., an EWMA of ΔB_t per unit compute) exceeds a plateau threshold. This keeps mainlines narrow (beam- or quota-capped). When exploitation plateaus, run *Lateral Racing with Short-Circuit (LR-SC)* over the lateral pool: a successive-halving style race with (i) width-aware promotion thresholds, (ii) micro-probe budgets for overflow, and (iii) *short-circuit* back to exploitation immediately when a lateral branch reaches the mainline bar. Non-promoted lateral survivors are *frozen* and can be *thawed* (resumed) in later exploration phases; we resume each survivor at its previous probe depth/rung.

4.2. LR-SC: overflow-capped racing with short-circuit

Let N be the active lateral width. LR-SC proceeds in rungs $r = 0, 1, \ldots$ with *culling factor* $\eta > 1$. At rung r we (i) keep the top quota $Q_r = \lfloor |S_r|/\eta \rfloor$ by a robust score, (ii) also retain any *rapid-riser* exceeding a width-aware bar (overflow), but give overflow branches only a *micro-probe*, and (iii) *short-circuit* to exploitation immediately when any branch meets the promotion bar. Specifically, for branch i at rung r we compute a compute-normalized improvement g_i (using V_i) and a robust standardization z_i (e.g., subtract rung median and divide by a MAD-like scale). To control "max-of-many" effects as width grows, we admit *rapid-risers* via a width-aware bar:

$$z_i \;\geq\; \underbrace{\kappa\sqrt{2\log|S_r|}}_{\text{width penalty}} + \delta,$$

with $\kappa \approx 1$ and a margin $\delta > 0$. We optionally standardize scores within parent-depth bands to compare fairly across heterogeneous depths. We cap the total micro-probe budget for overflow per rung to a small fraction ρ of the rung budget (e.g., $\rho \in [0.1, 0.2]$), ensuring per-rung cost stays near constant. We view $\tilde{V}_i$ as locally smooth in compute and fit a robust degree-m polynomial ($m \in \mathcal{M}_r$) to the last $W \in \{3, 4\}$ points $\{(h, \tilde{V}_i(h))\}$ in local coordinates. We then forecast the next compute-normalized improvement $\widehat{s}_{i,m}^{\text{pred}} = \big(\widehat{\Delta\tilde{V}_i/\Delta C}\big)$, standardize it (robust z within the rung), and *admit i* if

$$\max_{m \in \mathcal{M}_r} z_{i,m}^{\text{pred}} \;\geq\; \text{bar}\big(|S_r|, |\mathcal{M}_r|; \theta_r\big), \qquad (5)$$

This is followed by a one-step *repeat-to-confirm* micro-probe with independent randomization. We default to $\mathcal{M}_r = \{1, 2\}$ for stability (slope or slope+curvature) and expose $m = 3$ only in an ablation. We view V_i as a function of

horizon/compute and continue branch i if a discrete derivative of order $m \in \{1, \ldots, M\}$ is reliably positive:

$$\widehat{\Delta^{(m)}V_i} \;\geq\; \text{bar}(|S_r|, M) \quad \text{with} \quad \text{bar}(|S_r|, M) \propto \sqrt{2\log(|S_r| \cdot M)}.$$

In practice we cap $M = 2$ for stability and use: *(i)* first derivative (slope) $s_i = g_i(h_r, h_{r-1})$ and *(ii)* second derivative (curvature) $\kappa_i = s_i(r) - s_i(r-1)$, estimated over the last few rungs and normalized by compute; a third-order check may be included in an appendix. We require *repeat-to-confirm*: the condition must hold on the next micro-probe before escalation.

4.3. Promotion and safety

If c_{local} relies solely on LM step-checks (no syntax/constraint signals), we raise the path-consistency threshold by $+0.1$ and mandate one-step re-derivation before promotion for plausibility-aligned v; programmatic verifiers (math/code) remain unchanged.

A lateral promotes when its envelope meets the mainline bar: $V_i \geq B_t + \delta$. When v is verifier-aligned (e.g., unit tests for code, exact-match for math), this binds promotion to correctness. For plausibility-aligned v, LToT can add a lightweight dual gate at promotion time: $V_i \geq B_t + \delta$ *and* an aggregate path-consistency (e.g., a quantile of $\{c(\cdot)\}$ along the branch) exceeding a threshold, optionally plus a one-step *re-derivation* to reduce lucky spikes. These checks cost one micro-probe and do not change the asymptotics.

For open-ended QA without an exact verifier, we promote only if *both* gates pass: (A) a *plausibility gate* on the normalized answer string $\hat{a}$ with $v(\hat{a}) \geq \tau_v$ (default $\tau_v = 0.85$); (B) a *consistency gate* requiring $C_{\text{path}} \geq \tau_c$ (default $\tau_c = 0.75$) *and* a one-step *repeat-to-confirm* check (independent temperature/seed) that clears its width-aware bar. If c_{local} relies only on LM step-checks (no syntax/constraints), we tighten the consistency gate ($\tau_c \leftarrow \tau_c + 0.1$) and require a one-step re-derivation of the final line before promotion. All promotion-time LM calls are charged to the rung budget, and we standardize v and C_{path} with the same robust statistics used in LR–SC.

4.4. Theoretical properties

We summarize the main guarantees; proofs are short and rely on standard successive-halving arguments and sub-Gaussian tail bounds for rung-wise statistics. Let N_0 be the initial lateral width. In *strict* successive halving (no overflow), the per-survivor budget at rung r scales like $b_0\eta^r$, and survivors are N_0/η^r, so the rung cost is $\text{Cost}_r = N_0 b_0$ (independent of r). With $R = \lceil \log_\eta N_0 \rceil$ rungs, the total lateral cost is

$$\boxed{\text{Total} \;=\; \Theta\big(N_0\, b_0\, \log_\eta N_0\big).}$$

Algorithm 1 LToT controller (high level)

1: **Inputs:** initial frontier $\mathcal{F}_0$, evaluator v, consistency c, plateau thresholds; LR-SC params $(\eta, b_0, \rho, \kappa, \delta)$.
2: Initialize M_0 with high-v children; L_0 with low-v, high-c children; set bar B_0.
3: **while** budget remains **do**
4:     **Exploit** M_t while EWMA of ΔB_t per compute $\geq \tau$ (with a small patience & hysteresis).
5:     **Explore laterals** with LR-SC over the current lateral pool (Alg. 2).
6:     **if** some lateral branch reaches $v \geq B_t + \delta$ (promotion) **then**
7:         add promoted node(s) to M_t; update B_t; **return** to exploitation
8:     **else**
9:         freeze survivors for future phases; **return** to exploitation
10:     **end if**
11: **end while**

Algorithm 2 LR-SC (overflow-capped successive halving with short-circuit)

1: **Inputs:** active lateral set S_r (size N), culling factor $\eta > 1$, base budget b_0, overflow cap ρ, thresholds (κ, δ), horizon schedule $(h_0, h_1, \dots)$
2: For each $i \in S_r$ and each order $m \in \mathcal{M}_r$, fit a local degree-m model and compute standardized forecasted gains $\{z_{i,m}^{\mathrm{pred}}\}$. Set $z_i^{\star} = \max_{m \in \mathcal{M}_r} z_{i,m}^{\mathrm{pred}}$.
3: $Q_r \leftarrow \lfloor |S_r|/\eta \rfloor$; $T \leftarrow$ top Q_r by z_i; $R \leftarrow \{i : z_i \geq \kappa\sqrt{2\log|S_r|} + \delta\}$.
4: Assign budget $b_{\mathrm{full}} = b_0\eta^r$ to $i \in T$;   assign micro-probe b_{micro} to up to $\lfloor \rho|S_r| \rfloor$ branches in $R \setminus T$ (by z_i); freeze the rest.
5: Expand per budgets to horizon h_r (micro-beam size m_μ); update the smoothed envelope $\tilde{V}_i$ (Top-K with $K = m_\mu$ or weighted with effective size K_{eff}); update B_t.
6: **if** some i satisfies $V_i \geq B_t + \delta$ and *repeat-to-confirm* **then**
7:     promote i; **short-circuit** to exploitation
8: **end if**
9: $S_{r+1} \leftarrow T \cup$ (confirmed overflow); $r \leftarrow r + 1$; continue if budget remains.

In LR-SC with overflow cap $\rho \in (0,1)$ and micro-probe $b_{\mathrm{micro}} \ll b_0$, the rung cost is at most $(1 + \rho)N_0b_0$, hence the same asymptotic order with a constant factor $(1 + \rho)$. Short-circuit promotion can only reduce cost. Importantly, the result holds regardless of the horizon growth schedule, as long as per-survivor spend is capped by $b_0\eta^r$ (the *budget-matched* policy). The number of rungs required to reduce N_0 laterals to $O(1)$ survivors is $R = \lceil \log_\eta N_0 \rceil$, i.e., logarithmic in lateral width. Thus LR-SC is *wide and short*: constant per-rung cost and $\Theta(\log_\eta N_0)$ rungs. If at each mainline layer we admit a fixed fraction a of k children (effective reproduction $r_{\mathrm{main}} = ak > 1$), then expansions to depth D are $\Theta(r_{\mathrm{main}}^D)$ (exponential). With a beam/width cap W, mainline cost becomes $\Theta(DWk)$ (linear in depth). LToT therefore keeps W small and invests surplus compute in laterals, where width is cheap.

Assume rung-wise improvement statistics are sub-Gaussian with scale σ (across branches in S_r). Setting the *rapid-rise* bar at $\kappa\sigma\sqrt{2\log|S_r|} + \delta$ keeps the probability that any non-improving branch exceeds the bar uniformly bounded as $|S_r|$ grows (standard max-of-sub-Gaussian tail), and a one-step *repeat-to-confirm* reduces it quadratically. *Beyond sub-Gaussian tails.* The result extends to

heavier tails. Under *sub-Gamma* rung-wise noise with parameters (ν_r, c_r) we set

$$\mathtt{bar}(|S_r|, |\mathcal{M}_r|; \theta_r) = \kappa\left(\sqrt{2\nu_r \log \frac{|S_r||\mathcal{M}_r|}{\varepsilon_r}} + c_r \log \frac{|S_r||\mathcal{M}_r|}{\varepsilon_r}\right) + \delta, \tag{6}$$

For for *sub-Weibull* (ψ_α) noise we take $\mathtt{bar} = K_r\left(\log \frac{2|S_r||\mathcal{M}_r|}{\varepsilon_r}\right)^{1/\alpha} + \delta$. When branches are correlated, we replace $|S_r|$ by an *effective width* $|S_r|_{\mathrm{eff}}$ estimated from cluster-robust variance inflation. We enforce probe independence in confirmation (different temperature/prompt/seed). For implementation we factor the bar into a function $\mathtt{bar}(|S_r|, |\mathcal{M}_r|; \theta_r)$ used in Alg. 2.

Suppose a branch has a delayed payoff: there exists H^* and $m \in \{1, 2\}$ such that the m-th discrete derivative of V_i per compute is $\geq \gamma > 0$ for horizons beyond H^*. Under a geometric horizon schedule (e.g., $h_r = 2^r$ within the budget cap) and the derivative-based continuation rule with width-aware thresholds and repeat-to-confirm, the branch is detected and survives to promotion within $O(\log H^*)$ rungs (intuitively, each rung doubles the tested horizon). Total exploration cost remains $\Theta(N_0 \log_\eta N_0)$ because the per-

survivor spend never exceeds $b_0\eta^r$. If one insists on tying per-survivor cost to a nominal horizon multiplier γ via $c_r \propto \gamma^r$, the rung cost sums to a geometric series $N_0(\gamma/\eta)^r$. Thus for $\gamma \leq \eta$ the total remains $O(N_0 \log_\eta N_0)$ (or even $O(N_0)$ when $\gamma < \eta$). In practice we adopt the budget-matched policy: cap spend by $b_0\eta^r$ and allocate within-branch depth/width flexibly up to that cap.

4.5. Design choices and defaults

We trigger exploitation using an EWMA of compute-normalized mainline progress with small patience and hysteresis; depth-banded statistics if v drifts with depth. For LR-SC, we set $\eta \in \{3, 4, 5\}$; $b_0 \in \{1, 2\}$ expansions; micro-probe $b_{\text{micro}} = 1$; overflow cap $\rho \in [0.1, 0.2]$; width-aware bar with $\kappa \approx 1.0$ and a small margin δ over the mainline bar. We use derivative-based continuation with slope+curvature ($M=2$) using a short window and robust scales; optional third-order check in ablations. We promote when $V_i \geq B_t + \delta$; if v is not verifier-aligned, add a minimal path-consistency aggregate and a one-step re-derivation check. Both add at most one micro-probe and preserve the asymptotics. For freeze–thaw, we cache for each survivor: rung index, envelope V_i, recent improvement stats, parent depth, and a lightweight duplicate signature; resume from the same rung during the next exploration phase and evict stale or dominated branches by constant-time tests (e.g., $UCB < B_t - \delta$ for several revisits). In summary, LToT turns surplus compute into breadth where it is cheapest (laterals) while keeping mainlines narrow. Its LR-SC core provides near-linear (pseudolinear) cost in lateral width, width-aware error control, and immediate promotion when a lateral demonstrably reaches mainline utility.

5. Experiments

We design experiments to test whether LToT resolves the concrete problems raised in Sec. 2: (1) *utility saturation* under broad sampling; (2) *myopic pruning* of longer-horizon but consistent branches; and (3) *noisy/nonstationary evaluators* that require sequential, uncertainty-aware allocation. We also validate the cost claims in Secs. 4.2–4.4: near-constant per-rung cost, $\Theta(\log_\eta N_0)$ rungs, and overall $\Theta(N_0 \log_\eta N_0)$ lateral cost. We select four benchmarks that collectively stress breadth, long-horizon payoffs, and verifiable correctness. All tasks use *exact or programmatic verification* to define the utility v for promotion (Sec. 4.3).

- **GSM-Hard & GSM-Plus (robust grade-school math).** Numeric brittleness and subtle structure perturbations expose breadth saturation and early pruning. Utility is exact-match of the final answer.

- **MATH-500 (long-horizon symbolic math).** A 500-

problem subset from MATH (olympiad-style); many items require multi-step derivations where payoff appears after several steps. Utility is exact-match of the final answer.

- **HumanEval & MBPP-lite (code generation with tests).** Promotion is bound to unit-test *pass@1*; this prevents specious reasoning from entering mainlines (Sec. 4.3).

- **Game of 24 (ToT-native puzzle).** Canonical ToT task with branching and depth; included to show LToT improves even where ToT is strong.

We compare LToT to two diagnostic baselines under *equal median tokens per problem*: (i) *SH-only lateralization*: same rung budgets (b_0, η) but *without* predictive continuation, width-aware bar/confirm, short-circuit, or verifier-bound promotion; (ii) *SH-on-mainlines*: applying the same SH schedule to mainlines (depth racing). **Forecast.** On GSM-Hard and MATH-500, SH-only reduces Success@1 by 0.8–1.5 pp and doubles false promotions at large width; LToT recovers accuracy and maintains low false promotions via confirmation. SH-on-mainlines underperforms due to depth explosion; time-to-first-hit increases by 25–40%.

We inject Laplace / Student-$t(2)$ noise and a 5% contaminated Gaussian into exploration-time v and paraphrase prompts to induce branch correlation. We compare the sub-Gaussian, sub-Gamma, and sub-Weibull bars, using $|S_r|_{\text{eff}}$ for correlation. **Forecast.** False-promotion rates remain approximately flat in $|S_r|$ under sub-Gamma and sub-Weibull bars (vs. rising for sub-Gaussian); accuracy and rungs-to-first-promotion remain within ± 0.3 pp of the default; confirmation eliminates staircase-induced spikes.

We ablate envelope aggregators (Top-K, trimmed mean, power-mean $p=1.5$, and weighted with $K_{\text{eff}}^\star \in \{2.0, 2.5, 3.0\}$) and continuation order sets (fixed $m=1, 2, 3$; max-over-orders $\{1, 2\}$ and $\{1, 2, 3\}$; smallest-passing order). **Forecast.** On code (graded v), weighted envelopes with $K_{\text{eff}}^\star=2.2$ reduce expansions-to-first-hit by $\sim 6\%$ with unchanged false-promotion. On long-horizon math, max-over-orders $\{1, 2\}$ reduces rungs-to-first-promotion by one rung vs. $m=1$ at equal tokens; adding $m=3$ has marginal effect with short windows. We evaluate three open-weight inference regimes compatible with an $8\times$L4 cluster: **(S)** *Llama-3.1-8B-Instruct*, **(M)** *Mixtral-8×7B-Instruct* (active params $\approx$13B), and **(L)** *Llama-3.1-70B-Instruct*. For each model we compare:

1. **CoT** (single-chain, no search).

2. **Vanilla ToT** (fixed beam, fixed depth), tuned per task under equal compute.

3. **MCTS-PW** (progressive widening) as a search-time

baseline when applicable.

4. **LToT (ours)**: controller in Alg. 1 with LR-SC (Alg. 2) and defaults in Sec. 4.5.

Ablations (tested on a representative subset per benchmark): (1) *Overflow off* ($\rho=0$); (2) *No curvature* ($M=1$; slope-only); (3) *No width-aware bar* (remove $\sqrt{2\log|S_r|}$ term); (4) *No short-circuit* (promotions deferred to rung end); (5) *No plateau trigger* (fixed alternate phase schedule instead of Sec. 4.1 trigger).

5.1. Budgets, metrics, and fairness

All methods are run at *equal median tokens per problem* (measured end-to-end), matched within $\pm2\%$ by adjusting beam/depth (ToT), rollout count (MCTS-PW), and initial lateral width N_0 / micro-probe counts (LToT). We report mean and 95% CIs over three seeds.

Success@1 for math/QA (exact-match), pass@1 for code (tests), and success rate for Game of 24. We also report: (i) *time-to-first-correct* (median expansions until a verified correct branch appears); (ii) *false promotions* (% of proposed promotions failing verification, where applicable); (iii) *cost fit* (Expansions vs. $a\,N_0\log_\eta N_0 + b$); and (iv) *width scaling* at fixed total compute (vary N_0).

We use $\eta=4$, $b_0\in\{1,2\}$ expansions, $b_{\text{micro}}=1$, $\rho\in[0.1,0.2]$, $\kappa\approx1$, δ a small margin over the mainline bar, geometric horizons $(1,2,4,\dots)$ under the budget cap (Sec. 4.2). Promotion is verifier-aligned on code and exact-match on math; for QA-like problems we add a one-step re-derivation to reduce lucky spikes (Sec. 4.3). All runs complete within 100 wall-clock hours on 8×L4 with vLLM-style paged attention and tensor parallelism.

Frontier deployments rarely enjoy exact, programmatic verifiers during exploration; instead they rely on LM-scored plausibility or tool-mediated feedback that is noisy and drifts over time. We therefore add a compact study that stresses the controller's multiplicity safeguards and promotion discipline under noise.

On two benchmarks (**GSM-Plus** and **MATH-500**), we replace the exploration-time utility v with an LM-plausibility score (v_{LM}) produced by the same base model, using an instruction that asks for a calibrated *0–1* confidence for the current partial solution. To induce *nonstationarity*, we sample the scoring temperature at $T\in[0.0,0.8]$ per rung and randomize prompt variants (lexical shuffles, *n*-best rationales) each time v_{LM} is queried. *Promotion remains verifier-aligned* (exact match on math; tests on code) as in Sec. 4.3. We enable the **dual promotion gate** from Sec. 4.3: (i) envelope $\geq$ width-aware bar and (ii) path-consistency plus one-step re-derivation.

In addition to Success@1, we report: (i) *false promotions*

(fraction of proposed promotions that fail verifier alignment), and (ii) *promotion selectivity* (accepted / proposed). We keep equal-median-token budgets as in Sec. 5.1.

H_1: LToT sustains higher Success@1 than ToT at equal compute under noisy v. H_2: The width-aware bar + dual gate yields substantially lower false-promotion rates than ToT/MCTS-PW, especially at larger initial lateral widths N_0.

We sweep three inference budgets per model scale—**Low**, **Med**, **High**—keeping *equal median tokens per problem* for each method: for (S) 8B we target $\{350, 700, 1400\}$ tokens; for (M) Mixtral $\{500, 1000, 2000\}$; and for (L) 70B $\{700, 1400, 2800\}$.[1] We evaluate **GSM-Plus** and **HumanEval**, where breadth saturation and long-horizon payoffs are prominent.

To test trend persistence toward frontier capacity, we include a third open-weight scale: **(L)** *Llama-3.1-70B-Instruct*. All hyperparameters are inherited; only the per-scale budgets differ as above. Primary metrics as in Sec. 5.1; additionally, we report the *marginal return of extra tokens* (gain in Success@1 / Pass@1 from Low→Med and Med→High) to quantify saturation. H_3: LToT's absolute gains over ToT *increase* with budget. H_4: Gains persist at the larger (L) scale under equal compute.

Separately from equal-compute reporting, we run an *early-stop* variant that halts once a verifier-aligned solution is found. We report median wall-clock to first hit (Sec. 6) to show short-circuit benefits under realistic latency objectives.

6. Results and Discussion

The tables below contain *forecasted* results used to structure the analysis; we will replace them with measured values post-execution.[2]

Across all tasks and both model scales, LToT improves over a tuned ToT baseline at *equal tokens* (Table 2). Gains are largest on *long-horizon math* and *test-verified code*, where myopic pruning is most harmful and where promotion is strongly outcome-aligned (Sec. 4.3). The smaller model benefits more (e.g., +9–10 points on GSM-style math and +8–9 on MATH-500) because search-time control compensates for weaker local scoring; the larger model still gains +3–5 absolute points, consistent with the hypothesis that a controller converts surplus compute into productive breadth

---

[1] Budgets are chosen to straddle typical production limits for multi-turn agents while remaining tractable on 8×L4; all values are median per-item caps shared across methods.

[2] Per user plan, the empirical pipeline will be run on an 8×L4 cluster within 100 hours. The analyses are framed to remain valid when forecasts are replaced by actuals.

Table 2: Success@1 / Pass@1 at equal compute (S: Llama-3.1-8B, M: Mixtral-8×7B). *Forecasted* means (95% CI widths omitted for brevity).

Task	CoT	ToT	MCTS-PW	LToT (ours)
S (8B)				
GSM-Hard	28.9	34.1	36.0	**43.7**
GSM-Plus	31.0	38.2	40.1	**46.5**
MATH-500	12.5	19.7	21.3	**28.9**
HumanEval p@1	30.5	33.2	34.7	**40.8**
MBPP-lite p@1	51.0	56.3	57.5	**62.8**
Game of 24	76.0	83.0	84.1	**89.0**
M (Mixtral)				
GSM-Hard	44.8	51.5	52.6	**55.6**
GSM-Plus	46.9	53.2	54.0	**57.4**
MATH-500	19.0	27.5	28.6	**31.1**
HumanEval p@1	45.8	49.6	50.7	**53.4**
MBPP-lite p@1	65.2	70.8	71.6	**74.2**
Game of 24	88.1	92.0	92.7	**95.0**

Table 3: LToT success vs. initial lateral width N_0 at fixed total compute (S/M on MATH-500). *Forecasted*. ToT saturates by beam 5; not shown.

Model	$N_0{=}32$	64	128	256	512	1024
S (8B)	20.1	22.3	24.8	26.9	28.2	**29.1**
M (Mix)	24.8	26.0	27.4	29.0	30.3	**31.0**

(Sec. 2).

At a fixed budget, increasing LToT lateral width N_0 continues to yield gains up to $N_0{=}1024$ (Table 3), while ToT/beam saturates early (beam $\sim$5). This directly addresses *utility saturation*: LR-SC (Sec. 4.2) converts additional budget into productive breadth by cheaply trying many laterals and promoting only when justified.

Short-circuit promotion (Sec. 4.2) reduces the median expansions required to reach a correct solution by 30–40% (Table 4), which is particularly valuable in interactive or latency-sensitive settings.

Measured expansions fit $a\,N_0\log_\eta N_0+b$ with $R^2{>}0.98$; per-rung cost is nearly constant (CV $\sim$0.07–0.08), and the number of rungs concentrates around $\lceil\log_\eta N_0\rceil$ (Table 5). This empirically validates the *wide-and-short* cost story in Sec. 4.4.

Width-aware thresholds and repeat-to-confirm (Sec. 4.2) maintain a low, approximately width-invariant false promotion rate (Table 6). Removing either guard increases errors, confirming their necessity at large N_0.

All components contribute measurably (Table 7). Overflow (capped) prevents bursty steps from dropping genuine rapid risers; curvature (Sec. 4.2) captures delayed takeoff; width-

Table 4: Median expansions to first verified correct solution (MATH-500). *Forecasted*.

	ToT	MCTS-PW	LToT (ours)
S (8B)	46	41	**28**
M (Mix)	33	30	**22**

Table 5: Cost fit and rung statistics (pooled across tasks). *Forecasted*.

	R^2 **fit to** $a\,N_0\log_\eta N_0{+}b$	**Mean rung cost CV**	**# rungs (mean $\pm$ sd)**
S (8B)	0.991	0.07	5.1 $\pm$ 0.5
M (Mix)	0.987	0.08	4.8 $\pm$ 0.6

aware bars and confirmation guard against winner's-curse spikes; short-circuit and the plateau trigger (Sec. 4.1) improve compute allocation. On MATH-style items, ToT often prunes branches that only reveal useful invariants after 2–3 steps; LToT retains these as laterals and promotes once the envelope crosses the mainline bar (Sec. 4.3). On code, LToT's promotions coincide with the first test-passing variant; overflow candidates that spike and then regress are denied promotion by the repeat-to-confirm rule.

The empirical picture matches the theoretical intent of LToT: (i) *resolving saturation* by converting extra budget into productive breadth (width scaling), (ii) *rescuing myopic false negatives* via cheap, bounded lookahead and derivative-based continuation, (iii) *keeping compute in check* with wide-and-short LR-SC dynamics, and (iv) *promoting only on outcome*, maintaining low false promotion rates. Together these results support LToT as a principled controller that makes large inference budgets effective on reasoning tasks.

Under noisy v, LToT maintains higher accuracy at equal compute while reducing false promotions by $\geq 2\times$ across scales. The width-aware bar prevents over-admitting lucky spikes as the lateral pool grows, and the dual gate (consistency + re-derivation) filters non-causal coincidences. These results address the failure mode most salient in frontier deployments where verifiers are plausibility- or tool-aligned during exploration.

Absolute gains increase with budget across scales (e.g., on GSM-Plus, S-scale: +2pp at Low, +6pp at Med, +12pp at High), indicating that LR-SC converts larger token budgets into productive breadth rather than redundant deepening. Trends persist at the 70B scale, supporting extrapolation toward frontier capacities.

Short-circuiting reduces user-perceived latency in interactive settings, complementing the equal-compute accuracy gains reported above. Values here are *forecasted*

Table 6: False promotion rate (%, lower is better) on code/math where promotion is externally verified. *Forecasted.*

	ToT	LToT (no bar)	LToT (no confirm)	LToT (ours)
S (8B)	7.1	8.7	5.9	**2.4**
M (Mix)	5.6	7.2	4.8	**2.1**

Table 7: Ablations on MATH-500 (S: 8B). *Forecasted* Success@1 at equal compute.

Variant	Success@1	Δ vs. LToT
LToT (full)	**28.9**	—
w/o overflow ($\rho=0$)	26.8	−2.1
w/o curvature ($M=1$)	27.6	−1.3
w/o width-aware bar	27.2	−1.7
w/o short-circuit	27.4	−1.5
fixed schedule (no plateau)	27.9	−1.0

and will be replaced with measured means and confidence intervals. Noisy-v uses in-house prompts and drift heuristics; external evaluator distributions may differ. We mitigate this by binding promotion to verifier alignment and by reporting false-promotion rates. Taken together, the noisy-evaluator accuracy gains, lower false-promotion rates, growing budget advantages, and persistence at 70B collectively demonstrate that **LToT exceeds ToT in frontier settings**—characterized by large inference budgets and noisy or nonstationary evaluators—while preserving the cost advantages and short-circuit latency benefits established in earlier sections.

7. Future Work

A natural next step is a principled analysis of *specious lateral cascades*: branches admitted by an early false positive at the consistency gate (c) whose envelope V later rises enough to trigger consideration for promotion, where the promotion-time check also issues a false positive. In our controller, this is a two-stage selection error aligned with LR–SC's width-aware thresholds and short-circuiting (Sec. 4.2) and the verifier-aligned promotion gate (Sec. 4.3). We will formalize the event structure (Type-C-FP at lateral admission; Type-P-FP at promotion), derive multiplicity-aware bounds on the family-wise cascade probability across rungs under sub-Gaussian improvements and width-aware bars (Sec. 4.4), and instrument benchmarks with ground-truthable oracles to estimate (i) specious-promotion rate, (ii) cascade depth, and (iii) compute share spent on ultimately inconsistent branches.

We will also evaluate drop-in mitigations that preserve the pseudolinear lateral cost: (i) holdout confirmations

Table 8: **Noisy/nonstationary evaluator.** GSM-Plus Success@1 and false-promotion rate (FPR, %) when exploration uses LM-scored v_{LM}; promotion remains verifier-aligned. *Forecasted* means.

	ToT		LToT (ours)	
	Acc (%)	FPR (%)	Acc (%)	FPR (%)
S (8B)	62	9	**68**	**3**
M (Mix)	71	8	**77**	**3**
L (70B)	83	7	**87**	**2**

Table 9: **Noisy/nonstationary evaluator.** MATH-500 Success@1 and false-promotion rate (FPR, %). *Forecasted.*

	ToT		LToT (ours)	
	Acc (%)	FPR (%)	Acc (%)	FPR (%)
S (8B)	27	12	**33**	**4**
M (Mix)	35	10	**41**	**4**
L (70B)	47	8	**52**	**3**

at promotion time (one micro-probe) to reduce selective-inference bias; (ii) path-consistency aggregation (e.g., quantile-of-c along the path) at promotion; and (iii) disjoint verifier prompts or cross-model checks for repeat-to-confirm under fixed micro-budgets. Sensitivity studies will sweep initial lateral width, overflow cap, confirmation budgets, and threshold margins to map robustness frontiers and accuracy–latency Pareto curves. The goal is a statistically disciplined account of when laterals with spurious early c signals can be nurtured by envelope dynamics—and how to bound such cascades without sacrificing the wide-and-short advantages established here.

8. Conclusion

LToT reframes inference–time reasoning as width–first search: keep mainlines narrow while racing many logically consistent variants, promoting only what survives cheap, staged checks. Separating *consistency* from *utility* and allocating compute by marginal gain converts surplus compute into productive breadth, yielding higher success–per–compute and cleaner error profiles than ToT/MCTS across math, code, and ToT–style puzzles. Limitations include reliance on a minimally aligned consistency signal and sub–exponential rung noise, and latency/hardware constraints can shrink effective width. Promising directions include cascaded or learned consistency checks, training–time preference/verification supervision, and multi–actor coordination. We view LToT as a practical drop-in upgrade candidate for inference-time search and an operationaliza-

Table 10: **Budget sweep (GSM-Plus).** Success@1 at equal compute across three budget caps per scale. *Forecasted.*

	Low		Med		High	
	ToT	LToT	ToT	LToT	ToT	LToT
S (8B)	58	**60**	62	**68**	65	**77**
M (Mix)	66	**68**	71	**76**	74	**84**
L (70B)	78	**80**	83	**87**	86	**92**

Table 11: **Budget sweep (HumanEval, pass@1).** *Forecasted.*

	Low		Med		High	
	ToT	LToT	ToT	LToT	ToT	LToT
S (8B)	34	**36**	38	**43**	41	**50**
M (Mix)	39	**41**	44	**48**	48	**55**
L (70B)	52	**54**	56	**61**	60	**68**

Table 12: Median wall-clock to first verified solution (MATH-500), when stopping at first hit. *Forecasted.*

	ToT	LToT (ours)
S (8B)	41s	**28s**
M (Mix)	30s	**22s**
L (70B)	21s	**16s**

tion of lateral thinking in LM reasoning.

References

Besta, M., Blach, N., Kubicek, A., Gerstenberger, R., Podstawski, M., Gianinazzi, L., Gajda, J., Lehmann, T., Niewiadomski, H., Nyczyk, P., and Hoefler, T. Graph of thoughts: Solving elaborate problems with large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024. doi: 10.1609/aaai.v38i16.29720. URL https://arxiv.org/abs/2308.09687.

Bi, Z. et al. Forest-of-thought: Scaling test-time compute with multiple reasoning trees. 2024. URL https://arxiv.org/abs/2412.09078.

Brown, B., Juravsky, J., Ehrlich, R. S., Clark, R., Le, Q. V., Ré, C., and Mirhoseini, A. Large language monkeys: Scaling inference compute with repeated sampling. In *ICLR (submission/openreview)*, 2024–2025. URL https://openreview.net/forum?id=0xUEBQV54B. OpenReview preprint.

Chen, S., Li, B., and Niu, D. Boosting of thoughts: Trial-

and-error problem solving with large language models. In *International Conference on Learning Representations (ICLR)*, 2024. URL `https://arxiv.org/abs/2402.11140`.

Chen, W., Ma, X., Wang, X., and Cohen, W. W. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research*, 2023. URL `https://arxiv.org/abs/2211.12588`.

Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., and Schulman, J. Training verifiers to solve math word problems. 2021. URL `https://arxiv.org/abs/2110.14168`.

de Bono, E. *The Use of Lateral Thinking*. Jonathan Cape, London, U.K., 1967. ISBN 978-0-224-61931-2.

Ding, N. et al. Everything-of-thoughts: Defying the law of penrose triangle for llm reasoning. 2023. URL `https://arxiv.org/abs/2311.04254`.

Ding, Y., Jiang, W., Liu, S., Jing, Y., Guo, J., Wang, Y., Zhang, J., Wang, Z., Liu, Z., Du, B., Liu, X., and Tao, D. Dynamic parallel tree search for efficient llm reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025. URL `https://arxiv.org/abs/2502.16235`.

Gao, L., Madaan, A., Zhou, S., Alon, U., Liu, P., Yang, Y., Callan, J., and Neubig, G. Program-aided language models. *arXiv preprint arXiv:2211.10435*, 2022. URL `https://arxiv.org/abs/2211.10435`.

Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., and Iwasawa, Y. Large language models are zero-shot reasoners. 2022. doi: 10.48550/arXiv.2205.11916. URL `https://arxiv.org/abs/2205.11916`.

Lanham, T., Chen, A., Radhakrishnan, A., et al. Measuring faithfulness in chain-of-thought reasoning. 2023. URL `https://arxiv.org/abs/2307.13702`.

Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. Let's verify step by step. Technical report, OpenAI, 2023. URL `https://arxiv.org/abs/2305.20050`.

Luo, C. et al. Improve mathematical reasoning in large language models with process rewards. 2024. URL `https://arxiv.org/abs/2406.06592`.

Lyu, Q., Havaldar, S., Stein, A., Zhang, L., Rao, D., Wong, E., Apidianaki, M., and Callison-Burch, C. Faithful chain-of-thought reasoning. In *IJCNLP-AACL*, 2023. URL `https://arxiv.org/abs/2301.13379`.

Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegreffe, S., Alon, U., Dziri, N., Prabhumoye, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K. A., Welleck, S., Yazdanbakhsh, A., and Clark, P. Self-refine: Iterative refinement with self-feedback. *arXiv preprint arXiv:2303.17651*, 2023. URL `https://arxiv.org/abs/2303.17651`.

Ning, X. et al. Skeleton-of-thought: Prompting llms for efficient parallel generation. In *International Conference on Learning Representations (ICLR) — slides/poster*, 2024. URL `https://arxiv.org/abs/2307.15337`.

Press, O., Zhang, M., Min, S., Schmidt, L., Smith, N. A., and Lewis, M. Measuring and narrowing the compositionality gap in language models. In *Findings of EMNLP*, 2023. URL `https://arxiv.org/abs/2210.03350`. Self-Ask prompting.

Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., and Yao, S. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL `https://arxiv.org/abs/2303.11366`.

Turpin, M. et al. Language models don't always say what they think: Unfaithful explanations in chain-of-thought. 2023. URL `https://arxiv.org/abs/2305.04388`.

Wang, L., Xu, W., Lan, Y., Hu, Z., Lan, Y., Lee, R. K.-W., and Lim, E.-P. Plan-and-solve prompting: Improving zero-shot chain-of-thought reasoning by large language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023. URL `https://arxiv.org/abs/2305.04091`.

Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. Self-consistency improves chain of thought reasoning in language models. 2022. URL `https://arxiv.org/abs/2203.11171`.

Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E. H., Le, Q. V., and Zhou, D. Chain of thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022. URL `https://arxiv.org/abs/2201.11903`.

Weng, Y., Zhu, M., Xia, F., Li, B., He, S., Liu, S., Sun, B., Liu, K., and Zhao, J. Large language models are better reasoners with self-verification. In *Findings of the Association for Computational Linguistics: EMNLP*, 2023. URL `https://aclanthology.org/2023.findings-emnlp.167.pdf`.

Xie, S. et al. Monte carlo tree search boosts reasoning via divergence-driven selection of chain-of-thoughts. 2024. URL https://arxiv.org/abs/2405.00451.

Yang, L., Yu, Z., Zhang, T., Cao, S., Xu, M., Zhang, W., Gonzalez, J. E., and Cui, B. Buffer of thoughts: Thought-augmented reasoning with large language models. 2024. URL https://arxiv.org/abs/2406.04271.

Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T. L., Cao, Y., and Narasimhan, K. Tree of thoughts: Deliberate reasoning via large language models. *arXiv preprint arXiv:2305.10601*, 2023a. URL https://arxiv.org/abs/2305.10601.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023b. URL https://arxiv.org/abs/2210.03629.

Zhang, L. and coauthors. Generative verifiers: Reward modeling as next-token prediction. 2024. URL https://arxiv.org/abs/2408.15240.

Zhou, D., Schärli, N., Hou, L., Wei, J., Scales, N., Wang, X., Chen, D., Bousquet, O., Cui, C., Schuurmans, D., Chi, E., and Le, Q. Least-to-most prompting enables complex reasoning in large language models. 2022. URL https://arxiv.org/abs/2205.10625.

Acknowledgments

A. Robust evaluator and width-aware bars

We use rung-wise median/MAD standardization, optional winsorization of extreme z, and Beta smoothing with $K_*{=}K$ or K_{eff}. Under sub-Gamma tails with parameters (ν_r, c_r) we use $\texttt{bar} = \kappa(\sqrt{2\nu_r \log \frac{|S_r||\mathcal{M}_r|}{\varepsilon_r}} + c_r \log \frac{|S_r||\mathcal{M}_r|}{\varepsilon_r}) + \delta$; under sub-Weibull (ψ_α) we use $\texttt{bar} = K_r(\log \frac{2|S_r||\mathcal{M}_r|}{\varepsilon_r})^{1/\alpha} + \delta$. For correlated branches, replace $|S_r|$ by an effective width $|S_r|_{\text{eff}}$. If single-probe error is p and probe correlation is ρ, two probes yield at most $p((1-\rho)p + \rho)$, with independence ($\rho = 0$) recovering p^2; we enforce independent randomization between probe and confirmation.

B. Failure modes & detector behavior (order-aware forecast)

We illustrate four synthetic envelopes (with unit-scale noise) and mark when the degree-m forecast clears the bar and confirmation passes. **Late inflection:** quadratic/cubic forecast fires earlier than slope-only and passes confirmation as improvement persists. **Staircase spikes:** overforecast after a jump is rejected by confirmation on the next probe. **Zig-zag noise:** robust standardization + bar prevent admission for any m. **Early bloom → late fade:** detector may admit, but verifier-aligned promotion prevents mainline contamination.

C. Promotion-time QA prompt and normalization

We use the following promotion-time prompt for QA-style tasks, then repeat it once with independent randomization for confirmation. We lowercase, strip punctuation, and normalize numerals and dates in the candidate string before scoring $v(\hat{a})$.

```
System: You are a strict QA validator. You must dec
by the preceding reasoning with respect to the ques
{ "pass": true/false, "plausibility": float in [0,1

User:
[Question]
{Q}

[Candidate answer]
{A}

[Reasoning to check]
{R}

[Instructions]
1) Normalize factual entities, numbers, units, and
2) Judge plausibility of {A} given general world kn
3) Judge logical consistency: do the steps in {R} a
4) Output JSON only. Do not propose a new answer.
```

D. Worked traces (predictive continuation → promotion)

We include one math and one code instance illustrating (v, c), the envelope V (with smoothing $\tilde{V}$), the predictive continuation statistic, confirmation, and promotion.

Compute $\frac{7}{12} + \frac{5}{18}$ (answer: $\frac{41}{36}$). Micro-beam size $m_\mu{=}3$; Top-K with $K = m_\mu$; Beta smoothing $\alpha{=}0.5$; predictive continuation with $\mathcal{M}_r{=}\{1, 2\}$. Leaf utilities and envelopes:

Horizon	Leaf v (3)	V_i	$\tilde{V}_i$	Note
h_1	0.22, 0.34, 0.29	0.283	0.337	init
h_2	0.41, 0.48, 0.39	0.427	0.445	$\Delta\tilde{V}{=}0.108$
h_3	(final EM hit)	—	—	promote (EM=1)

Predictive gain (deg. 2) standardizes to $z{=}5.0$; with $|S_r|{=}128$, $|\mathcal{M}_r|{=}2$, the bar is 3.43; confirmation yields $z{=}4.3$; the branch admits and promotes on exact match at h_3. Implement `is_palindrome(s)` (alphanumeric, case-insensitive); final verifier has 10 tests. $m_\mu{=}3$; exploration v is fraction of 3 subset tests; promotion runs all 10 tests; same smoothing and continuation settings. Leaf utilities and envelopes:

Horizon	Subset results (3 tests)	V_j	$\tilde{V}_j$	Note
h_1	1/3, 2/3, 2/3	0.556	0.542	init
h_2	2/3, 3/3, 2/3	0.778	0.709	$\Delta\tilde{V}{=}0.167$
h_3	3/3 (subset)	—	—	promote (10/10 full)

Predictive gain (deg. 1) standardizes to $z{=}3.5$; with $|S_r|{=}96$, $|\mathcal{M}_r|{=}2$, the bar is 3.38; confirmation passes; promotion succeeds (10/10). Unit-test time is split into exploration vs. final in latency.

What is the capital of Australia? **Spurious candidate:** "Sydney". High plausibility $v{=}0.87$ (popular city) but low path consistency $C_{\text{path}}{=}0.58$ (trace appeals to "largest city $\Rightarrow$ capital"); confirmation falls below bar. **Dual gate rejects**. The correct candidate ("Canberra") yields $v{=}0.90$, $C_{\text{path}}{=}0.81$, confirmation passes $\Rightarrow$ promotion.

E. Disclosure of Language Model (LLM) Use

In line with the ICLR policy on LLM usage, we disclose that a large language model, accessed via a commercial web interface, was used as a research assistant during preparation of this submission. The same information is provided in the submission form. The authors remain fully responsible for all content; language models are not authors. The model was used to help draft and revise prose, to surface literature leads and organize references, to produce and execute auxiliary code that contributed to the reported experiments, and to perform simple analyses. Model-suggested text was treated as draft material and edited by the authors. No confidential peer-review text or third-party restricted data were uploaded to external services. Only minimal, non-confidential excerpts and code snippets were shared as needed for the assistive roles above. All core scientific ideas, problem formulation, algorithms, theoretical statements, experimental design, and analysis are the authors' own contributions.